



# Linear Filtering and Image Transforms

CSC 249/449 Spring 2019

<http://www.cs.rochester.edu/~cxu22/t/249S19/>

Instructor: Chenliang Xu  
chenliang.xu@rochester.edu

# Assignments and Python Tutorial


- Homework 1 is uploaded on Blackboard
- Due 1/31 midnight.
- We will have a Python tutorial Thu 1/24 7:30pm in 1400 Wegmans.

# Today's Content

- Images as functions
- Point operators
- Neighborhood operators
  - Linear filters
  - Nonlinear filters




(a)



(b)

45	60	98	127	132	133	137	133
46	65	98	123	126	128	131	133
47	65	96	115	119	123	135	137
47	63	91	107	113	122	138	134
50	59	80	97	110	123	133	134
49	53	68	83	97	113	128	133
50	50	58	70	84	102	116	126
50	50	52	58	69	86	101	120



(c)



(d)

**Figure 3.3** Visualizing image data: (a) original image; (b) cropped portion and scanline plot using an image inspection tool; (c) grid of numbers; (d) surface plot. For figures (c)–(d), the image was first converted to grayscale.

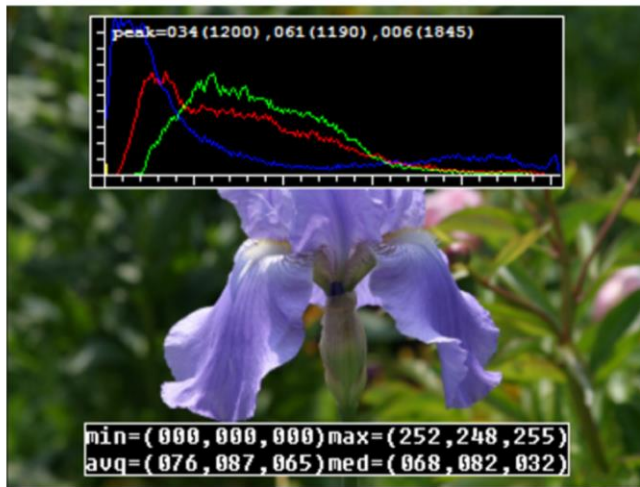
# Images as functions



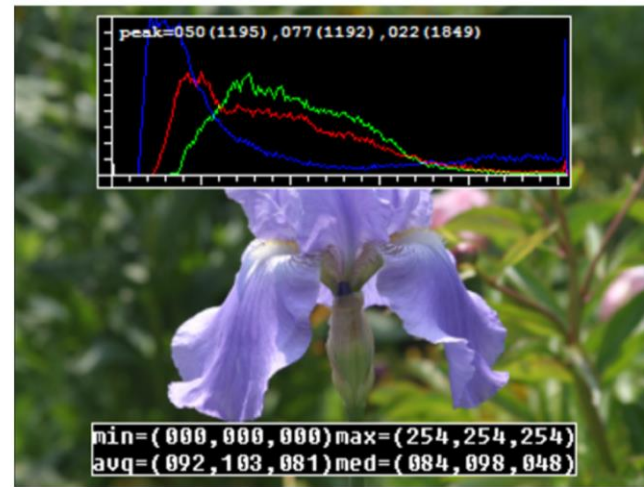
# Point Operators: Examples

Original Image

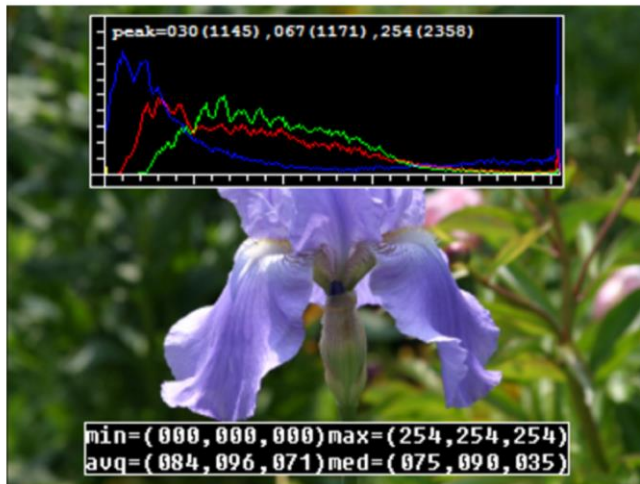
Brightness increased (additive offset,  $b=16$ )



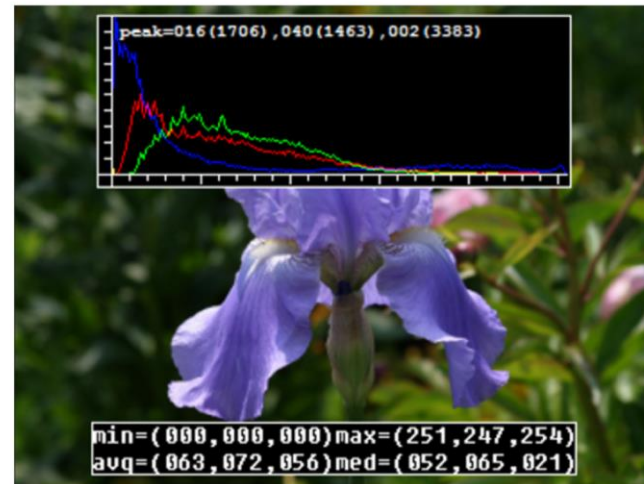
(a)



(b)



(c)





(d)

Contrast increased (multiplicative gain,  $a=1.1$ )      Gamma (partially) linearized ( $r=1.2$ )

# Point Operators: Examples

- Color transforms

RGB primaries



RGB to HSV color table


Color	Color name	Hex	(R,G,B)	(H,S,V)
	Black	#000000	(0,0,0)	(0°,0%,0%)
	White	#FFFFFF	(255,255,255)	(0°,0%,100%)
	Red	#FF0000	(255,0,0)	(0°,100%,100%)
	Lime	#00FF00	(0,255,0)	(120°,100%,100%)
	Blue	#0000FF	(0,0,255)	(240°,100%,100%)
	Yellow	#FFFF00	(255,255,0)	(60°,100%,100%)
	Cyan	#00FFFF	(0,255,255)	(180°,100%,100%)
	Magenta	#FF00FF	(255,0,255)	(300°,100%,100%)
	Silver	#C0C0C0	(192,192,192)	(0°,0%,75%)
	Gray	#808080	(128,128,128)	(0°,0%,50%)
	Maroon	#800000	(128,0,0)	(0°,100%,50%)
	Olive	#808000	(128,128,0)	(60°,100%,50%)
	Green	#008000	(0,128,0)	(120°,100%,50%)
	Purple	#800080	(128,0,128)	(300°,100%,50%)
	Teal	#008080	(0,128,128)	(180°,100%,50%)
	Navy	#000080	(0,0,128)	(240°,100%,50%)

# Defining convolution

- Let  $f$  be the image and  $g$  be the kernel. The output of convolving  $f$  with  $g$  is denoted  $f * g$ .

$$(f * g)[m, n] = \sum_{k, l} f[m - k, n - l] g[k, l]$$


Convention:  
kernel is “flipped”



- MATLAB functions: [conv2](#), [filter2](#), [imfilter](#)

# Dealing with edges

- What is the size of the output?
- MATLAB: `filter2(g, f, shape)`
  - *shape* = 'full': output size is sum of sizes of *f* and *g*
  - *shape* = 'same': output size is same as *f*
  - *shape* = 'valid': output size is difference of sizes of *f* and *g*





# Dealing with edges

- What about missing pixel values?
  - the filter window falls off the edge of the image
  - need to extrapolate
  - methods:
    - clip black (zero-padding)
    - wrap around
    - copy edge
    - reflect across edge



# Dealing with edges

- What about missing pixel values?
  - the filter window falls off the edge of the image
  - need to extrapolate
  - methods (MATLAB):
    - clip black (zero-padding): `imfilter(f, g, 0)`
    - wrap around: `imfilter(f, g, 'circular')`
    - copy edge: `imfilter(f, g, 'replicate')`
    - reflect across edge: `imfilter(f, g, 'symmetric')`

# Practice with linear filters



Original

0	0	0
0	1	0
0	0	0

?

# Practice with linear filters



Original

0	0	0
0	1	0
0	0	0



Filtered  
(no change)

# Practice with linear filters



Original

0	0	0
0	0	1
0	0	0

?

# Practice with linear filters



Original

0	0	0
0	0	1
0	0	0



Shifted *right*  
By 1 pixel

# Practice with linear filters



Original

 $\frac{1}{9}$ 

1	1	1
1	1	1
1	1	1

?

# Practice with linear filters



Original

$$\frac{1}{9}$$

1	1	1
1	1	1
1	1	1



Blur (with a  
box filter)



# Practice with linear filters



Original

0	0	0
0	2	0
0	0	0

—

$\frac{1}{9}$

1	1	1
1	1	1
1	1	1

?

(Note that filter sums to 1)

# Practice with linear filters



Original

0	0	0
0	2	0
0	0	0

—

$\frac{1}{9}$


1	1	1
1	1	1
1	1	1




## Sharpening filter

- Accentuates differences with local average

# Sharpening





**before**



**after**


# Smoothing with box filter revisited

- What's wrong with this picture?
- What's the solution?



# Smoothing with box filter revisited



- What's wrong with this picture?
- What's the solution?
  - To eliminate edge effects, weight contribution of neighborhood pixels according to their closeness to the center



“fuzzy blob”

# Gaussian Kernel

$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$




0.003	0.013	0.022	0.013	0.003
0.013	0.059	0.097	0.059	0.013
0.022	0.097	0.159	0.097	0.022
0.013	0.059	0.097	0.059	0.013
0.003	0.013	0.022	0.013	0.003

5 x 5,  $\sigma = 1$


- Constant factor at front makes volume sum to 1 (can be ignored when computing the filter values, as we should renormalize weights to sum to 1 in any case)

# Gaussian Kernel

$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$



$\sigma = 2$  with  $30 \times 30$   
kernel





$\sigma = 5$  with  $30 \times 30$   
kernel

- Standard deviation  $\sigma$ : determines extent of smoothing


# Choosing kernel width

- The Gaussian function has infinite support, but discrete filters use finite kernels





# Gaussian vs. box filtering



# Gaussian filters

- Remove high-frequency components from the image (*low-pass filter*)
- Convolution with self is another Gaussian
  - So can smooth with small- $\sigma$  kernel, repeat, and get same result as larger- $\sigma$  kernel would have
  - Convoluting two times with Gaussian kernel with std. dev.  $\sigma$  is same as convoluting once with kernel with std. dev.  $\sigma\sqrt{2}$
- *Separable* kernel
  - Factors into product of two 1D Gaussians
  - Discrete example:

$$\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \begin{bmatrix} 1 & 2 & 1 \end{bmatrix}$$

# Separability of the Gaussian filter

$$\begin{aligned} G_{\sigma}(x, y) &= \frac{1}{2\pi\sigma^2} \exp -\frac{x^2 + y^2}{2\sigma^2} \\ &= \left( \frac{1}{\sqrt{2\pi}\sigma} \exp -\frac{x^2}{2\sigma^2} \right) \left( \frac{1}{\sqrt{2\pi}\sigma} \exp -\frac{y^2}{2\sigma^2} \right) \end{aligned}$$

The 2D Gaussian can be expressed as the product of two functions, one a function of  $x$  and the other a function of  $y$

In this case, the two functions are the (identical) 1D Gaussian

# Non-linear Filters (and denoising)

# Noise



Original



Salt and pepper noise



Impulse noise




Gaussian noise

- **Salt and pepper noise:** contains random occurrences of black and white pixels
- **Impulse noise:** contains random occurrences of white pixels
- **Gaussian noise:** variations in intensity drawn from a Gaussian normal distribution

# Gaussian noise

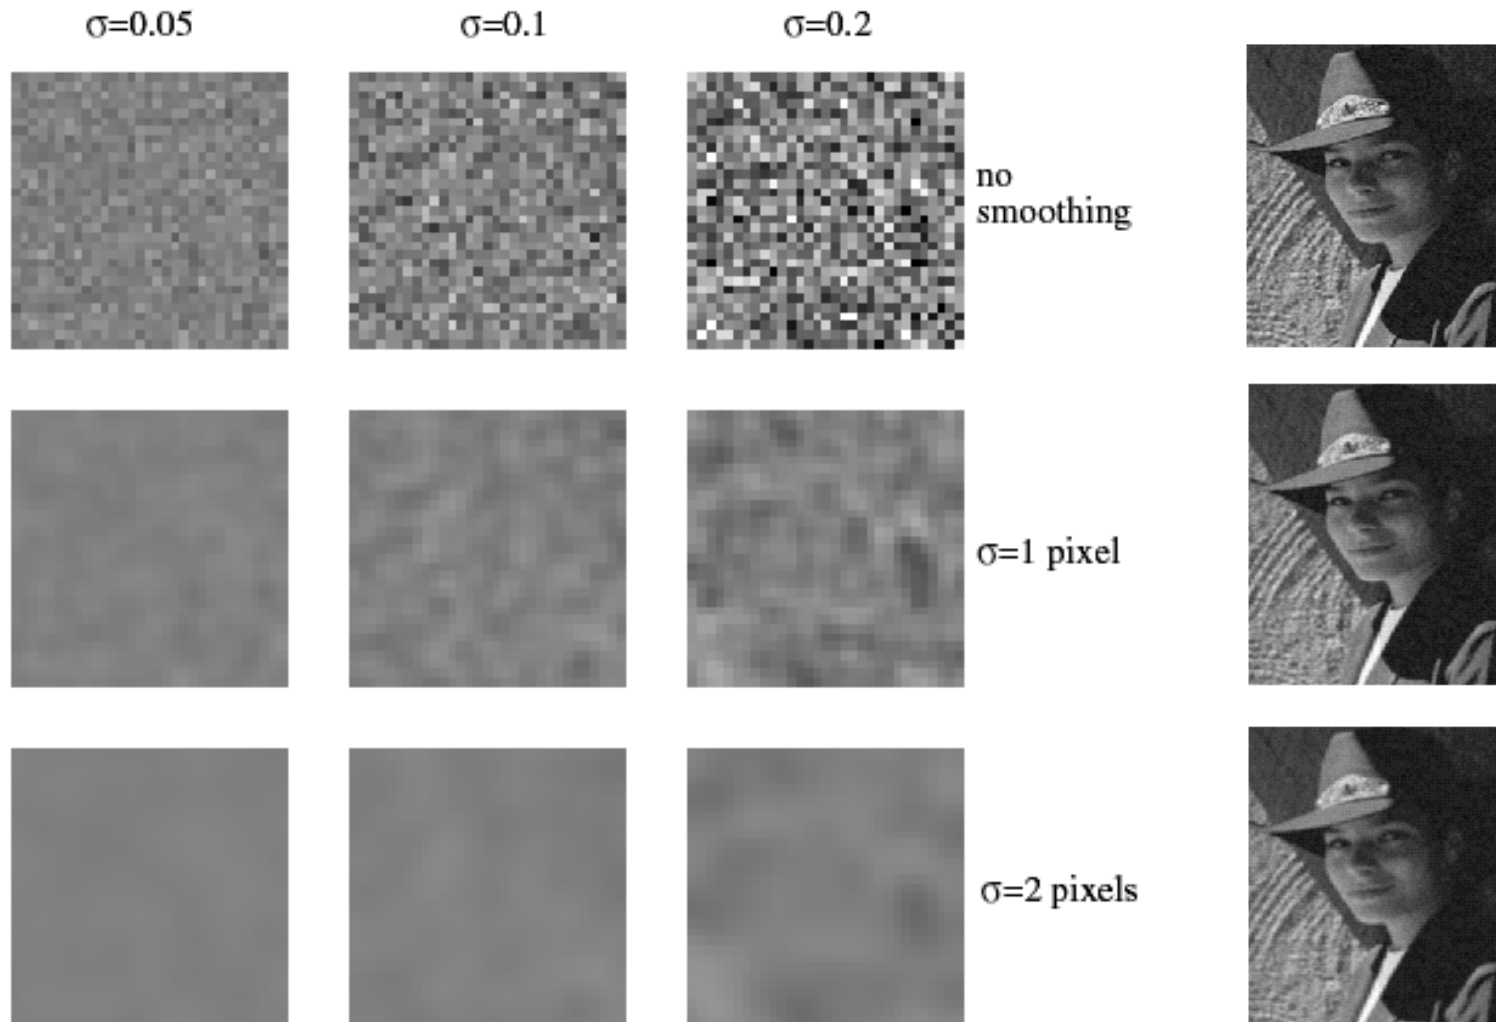
- Mathematical model: sum of many independent factors
- Good for small standard deviations
- Assumption: independent, zero-mean noise



$$f(x, y) = \overbrace{\hat{f}(x, y)}^{\text{Ideal Image}} + \overbrace{\eta(x, y)}^{\text{Noise process}}$$

Gaussian i.i.d. ("white") noise:  
 $\eta(x, y) \sim \mathcal{N}(\mu, \sigma)$

# Reducing Gaussian noise



Smoothing with larger standard deviations suppresses noise, but also blurs the image

# Reducing salt-and-pepper noise

3x3



5x5



7x7




- What's wrong with the results?



## Alternative idea: Median filtering

- A **median filter** operates over a window by selecting the median intensity in the window




- Is median filtering linear?

# Median filter



- What advantage does median filtering have over Gaussian filtering?
  - Robustness to outliers

filters have width 5 :




# Median filter

Salt-and-pepper noise



Median filtered



- MATLAB: `medfilt2(image, [h w])`

# Gaussian vs. median filtering

3x3

5x5

7x7


Gaussian




Median



# Sharpening revisited



before



after




# Sharpening revisited


- What does blurring take away?




-



=



Let's add it back:

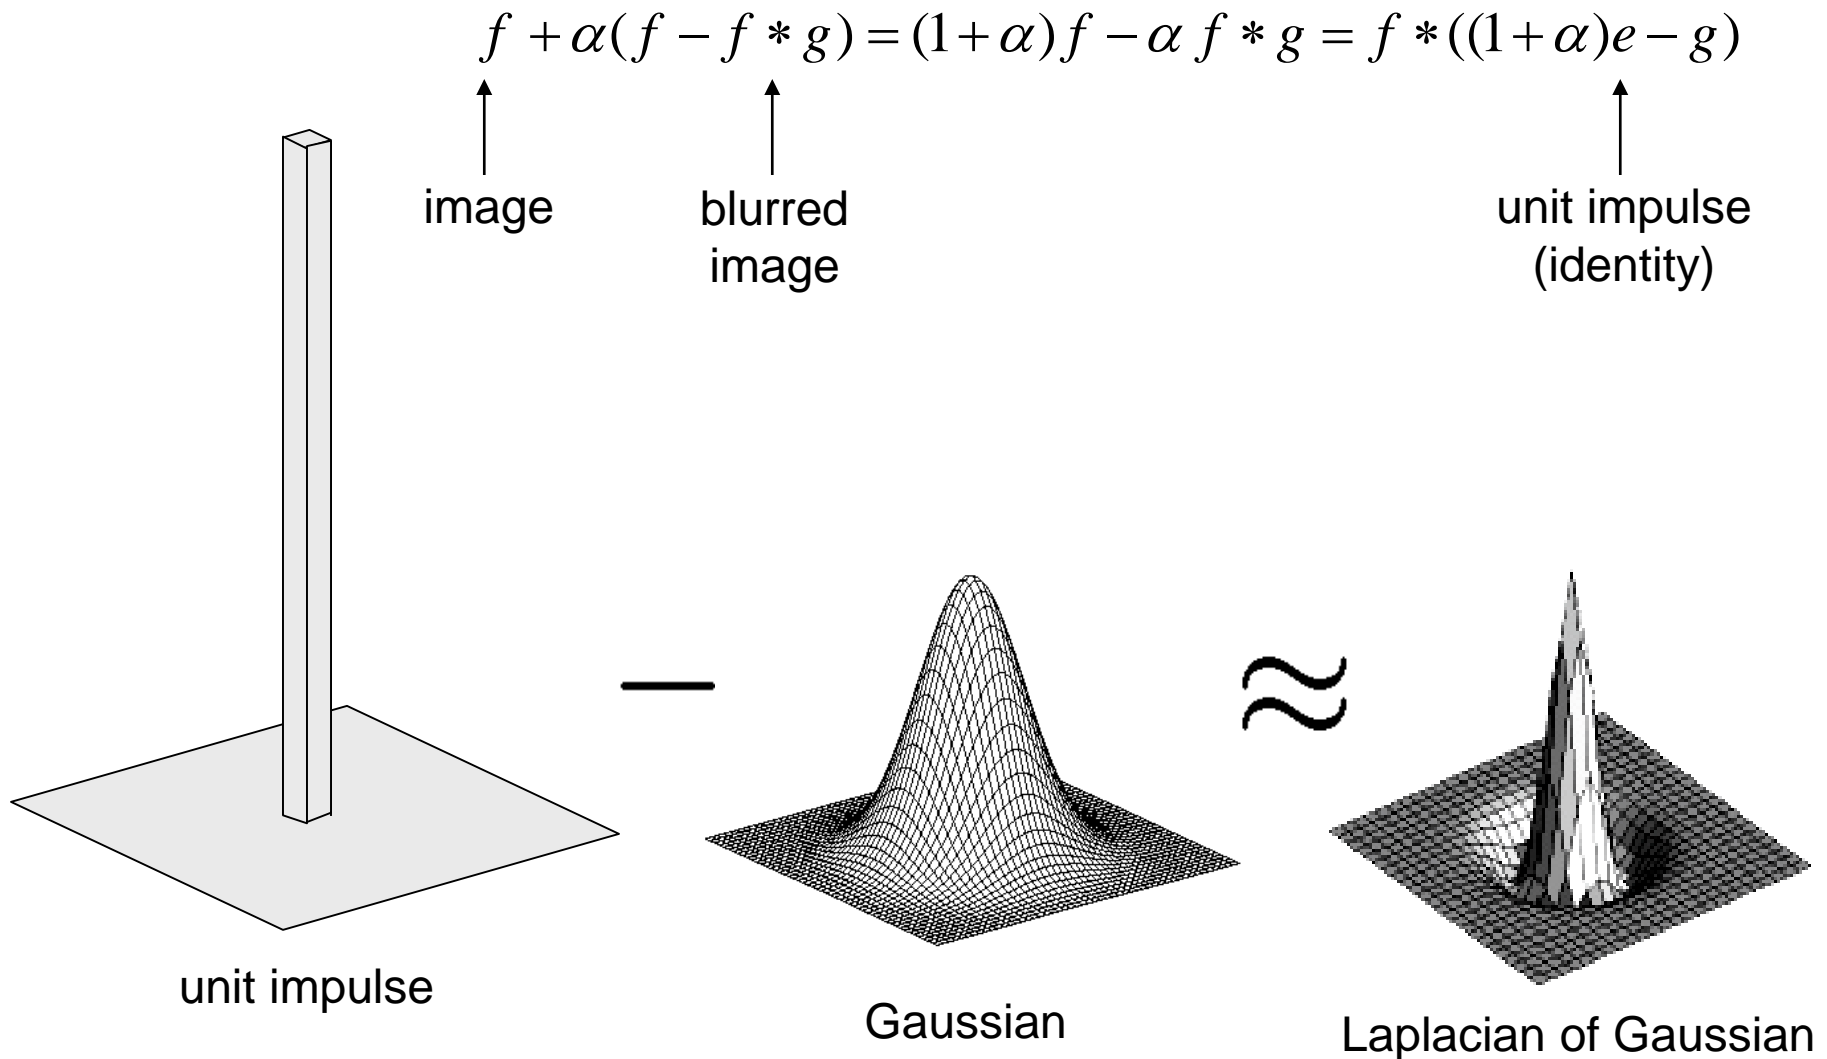
+  $\alpha$ 

=



Source: S. Lazebnik

# Unsharp mask filter



# Application: Hybrid Images



- A. Oliva, A. Torralba, P.G. Schyns, [“Hybrid Images,”](#) SIGGRAPH 2006



# Changing expression



Sad




Surprised



# Application: Hybrid Images

Gaussian Filter



Laplacian Filter

- A. Oliva, A. Torralba, P.G. Schyns, ["Hybrid Images,"](#) SIGGRAPH 2006

# Summary

- Images as functions
- Point operators
- Neighborhood operators
  - Linear filters
  - Nonlinear filters
- Optional Reading:
  - Fourier transforms (Images as points), Sec. 3.4.1-2
- Next Lecture:
  - Edges and Corners
  - Sec. 3.2.3, 4.2, 4.1.1 (Up to P. 190)